




# **Reinforcement Learning with Multiple Experts: A Bayesian Model Combination Approach**

(NIPS 2018)

Method : Replication

Team : 王如均, 黃立鈞, 潘柏瑋, 廖晔程

Date : 2022/6/13





# Outline

**01.** Introduction

**04.** Experimental Results


**02.** Related work

**05.** Conclusion

**03.** Bayesian Reward Shaping



# *01.* Introduction







# Problem

- Potential based reward shaping is a powerful technique for accelerating convergence of reinforcement learning algorithms. Such information includes an estimate of the optimal value function. However, this information is often biased or inaccurate and can mislead many reinforcement learning algorithms.

# Contribution

- This paper apply Bayesian Model Combination with multiple experts in a way that learns to trust a good combination of experts as training progresses. It is both computationally efficient and general, and is shown numerically to improve convergence across discrete and continuous domains.
- 

# *02.* Related work





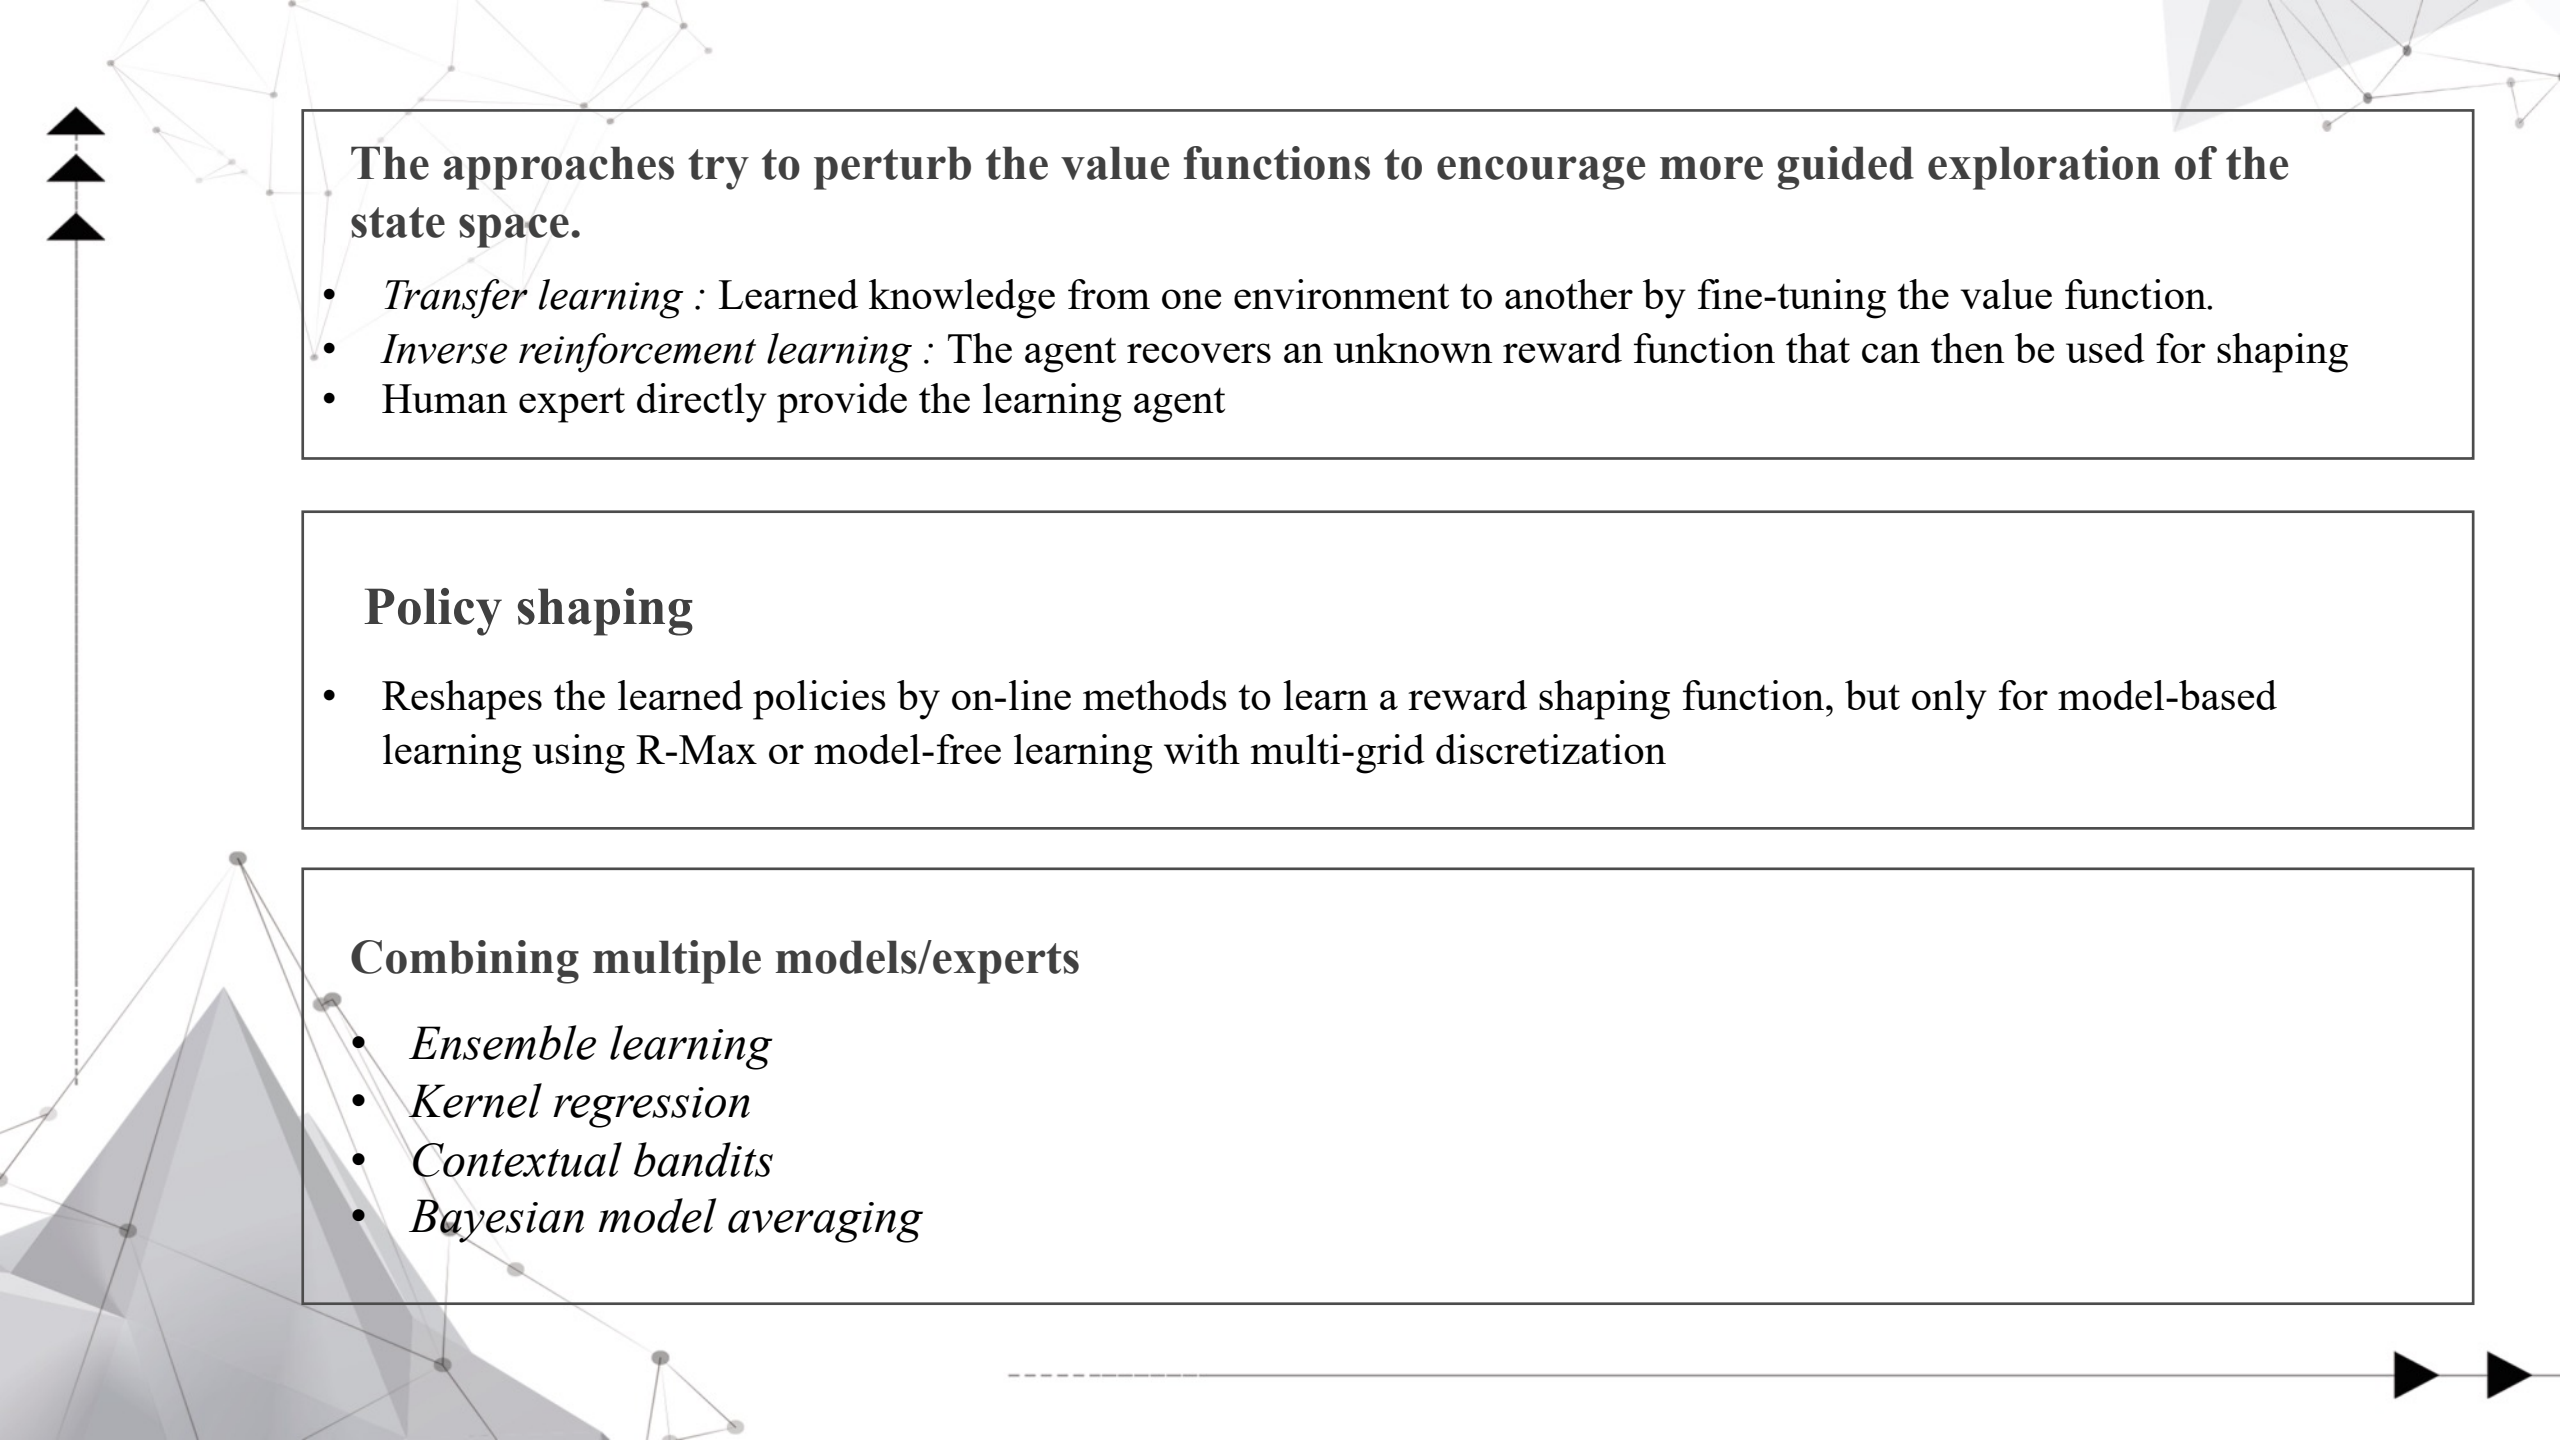
**The approaches try to perturb the value functions to encourage more guided exploration of the state space.**

- *Transfer learning* : Learned knowledge from one environment to another by fine-tuning the value function.
- *Inverse reinforcement learning* : The agent recovers an unknown reward function that can then be used for shaping
- Human expert directly provide the learning agent

## **Policy shaping**


- Reshapes the learned policies by on-line methods to learn a reward shaping function, but only for model-based learning using R-Max or model-free learning with multi-grid discretization

## **Combining multiple models/experts**

- *Ensemble learning*
  - *Kernel regression*
  - *Contextual bandits*
  - *Bayesian model averaging*
- 



# *03.* Bayesian Reward Shaping



## *Potential-Based Reward Shaping*

The idea of *reward shaping* is to incorporate prior knowledge about the domain in the form of additional rewards during training to speed up convergence towards the optimal policy

- Form of additional rewards

$$R'(s, a, s') = R(s, a, s') + F(s, a, s')$$

- *Potential-based reward shaping*

$$F(s, a, s') = \gamma\Phi(s') - \Phi(s)$$

- Non-stationary time-dependent potential functions

$$F(s, a, t, s', t') = \gamma\Phi(s', t') - \Phi(s, t)$$






## *Bayesian Reward Shaping*

<i>Bayesian model averaging (BMA)</i>	✓ <i>Bayesian model combination (BMC)</i>
BMA converges asymptotically toward the optimal <i>hypothesis</i>	BMC converges toward the optimal <i>ensemble</i>

The BMC approach has two clear advantages over BMA:

- (1) when two or more potential functions are optimal, it will converge to a linear combination of them
- (2) It provides an estimator with reduced variance.

This paper shows how BMC can be used to incorporate imperfect advice from multiple experts into RL problems, all with the same space and time complexity as TD-learning.



## Bayesian Model Combination

- $\mathbf{w}$  (weight vectors) : Interpreted as categorical distributions over experts and can learn the optimal distribution over experts

$$\mathcal{S}^{N-1} = \left\{ \mathbf{w} \in \mathbb{R}^N : \sum_{i=1}^N w_i = 1, w_i \geq 0 \right\}$$

$$\rho_t(s, a) = \mathbb{E}[q_{s,a} | \mathcal{D}] = \int_{\mathbb{R}} q \mathbb{P}(q | \mathcal{D}) \, dq$$

$$= \int_{\mathbb{R}} q \int_{\mathcal{S}^{N-1}} \mathbb{P}(q | \mathcal{D}, \mathbf{w}) \mathbb{P}(\mathbf{w} | \mathcal{D}) \, d\mathbf{w} \, dq = \int_{\mathbb{R}} q \int_{\mathcal{S}^{N-1}} \sum_{i=1}^N \mathbb{P}(q|i) w_i \pi_t(\mathbf{w}) \, d\mathbf{w} \, dq$$

$$= \sum_{i=1}^N \int_{\mathbb{R}} q \mathbb{P}(q|i) \int_{\mathcal{S}^{N-1}} w_i \pi_t(\mathbf{w}) \, d\mathbf{w} \, dq = \sum_{i=1}^N \int_{\mathbb{R}} q \mathbb{P}(q|i) \mathbb{E}_{\pi_t}[w_i] \, dq$$

$$= \sum_{i=1}^N \mathbb{E}_{\pi_t}[w_i] \int_{\mathbb{R}} q \mathbb{P}(q|i) \, dq = \sum_{i=1}^N \mathbb{E}_{\pi_t}[w_i] \mathbb{E}[q_{s,a}|i],$$

The total return can be written as a linear combination of individual return “contributions” from each expert model, weighted by the expected posterior belief that the expert is correct .

# Posterior Approximation using Moment Matching

Bayes' theorem

$$\pi_{t+1}(\mathbf{w}) = \mathbb{P}(\mathbf{w}|\mathcal{D}, d) \propto \mathbb{P}(d|\mathbf{w})\pi_t(\mathbf{w}) \propto \sum_{i=1}^N \mathbb{P}(d|i) \mathbb{P}(i|\mathbf{w})\pi_t(\mathbf{w}) = \frac{1}{C_{t+1}} \sum_{i=1}^N e_i w_i \pi_t(\mathbf{w}),$$

$$C_{t+1} = \int_{\mathcal{S}^{N-1}} \sum_{i=1}^N e_i w_i \pi_t(\mathbf{w}) d\mathbf{w} = \sum_{i=1}^N e_i \int_{\mathcal{S}^{N-1}} w_i \pi_t(\mathbf{w}) d\mathbf{w} = \sum_{i=1}^N e_i \mathbb{E}_{\pi_t}[w_i]$$

Multivariate *Dirichlet* distribution with parameters  $\alpha$

- Generalized moments

$$\mathbb{E}_f \left[ \prod_{i=1}^N w_i^{n_i} \right] = \frac{\Gamma \left( \sum_{i=1}^N \alpha_i \right)}{\Gamma \left( \sum_{i=1}^N (\alpha_i + n_i) \right)} \prod_{i=1}^N \frac{\Gamma(\alpha_i + n_i)}{\Gamma(\alpha_i)}, \quad n_i \geq 0.$$

- Density function

$$f(\mathbf{w}; \boldsymbol{\alpha}) = \frac{\Gamma \left( \sum_{i=1}^N \alpha_i \right)}{\prod_{i=1}^N \Gamma(\alpha_i)} \prod_{i=1}^N w_i^{\alpha_i - 1}, \quad \mathbf{w} \in \mathcal{S}^{N-1}$$

## Posterior Approximation using Moment Matching (Cont.)

- Approximate moment matching with proposal Dir ( $\alpha$ ) by solving the system of equations

$$m_i = \frac{\alpha_i}{\alpha_0}, \quad i = 1, 2 \dots N - 1$$

$$s_1 = \frac{\alpha_1(\alpha_1 + 1)}{\alpha_0(\alpha_0 + 1)}$$

Unique positive solution

$$\alpha_0 = \frac{m_1 - s_1}{s_1 - m_1^2}$$

$$\alpha_i = m_i \alpha_0 = m_i \left( \frac{m_1 - s_1}{s_1 - m_1^2} \right), \quad i = 1, 2 \dots N - 1$$

$$\begin{aligned} m_i &= \mathbb{E}_{\pi_{t+1}} [w_i] = \int_{\mathcal{S}^{N-1}} \frac{\alpha_{t,0}}{\mathbf{e} \cdot \boldsymbol{\alpha}_t} \sum_{j=1}^N e_j w_j w_i \pi_t(\mathbf{w}) \, d\mathbf{w} \\ &= \frac{\alpha_{t,0}}{\mathbf{e} \cdot \boldsymbol{\alpha}_t} \sum_{j=1}^N e_j \int_{\mathcal{S}^{N-1}} w_j w_i \pi_t(\mathbf{w}) \, d\mathbf{w} = \frac{\alpha_{t,0}}{\mathbf{e} \cdot \boldsymbol{\alpha}_t} \sum_{j=1}^N e_j \mathbb{E}_{\pi_t} [w_i w_j] \\ &= \frac{\alpha_{t,0}}{\mathbf{e} \cdot \boldsymbol{\alpha}_t} \left( e_i \mathbb{E}_{\pi_t} [w_i^2] + \sum_{j \neq i} e_j \mathbb{E}_{\pi_t} [w_i w_j] \right) \\ &= \frac{\alpha_{t,0}}{\mathbf{e} \cdot \boldsymbol{\alpha}_t} \left( e_i \frac{\alpha_{t,i}(\alpha_{t,i} + 1)}{\alpha_{t,0}(\alpha_{t,0} + 1)} + \sum_{j \neq i} e_j \frac{\alpha_{t,i} \alpha_{t,j}}{\alpha_{t,0}(\alpha_{t,0} + 1)} \right) \\ &= \frac{\alpha_{t,i}(e_i + \mathbf{e} \cdot \boldsymbol{\alpha}_t)}{(\mathbf{e} \cdot \boldsymbol{\alpha}_t)(\alpha_{t,0} + 1)}. \end{aligned}$$

$$s_1 = \frac{\alpha_{t,1}(\alpha_{t,1} + 1)(2e_1 + \mathbf{e} \cdot \boldsymbol{\alpha}_t)}{(\mathbf{e} \cdot \boldsymbol{\alpha}_t)(\alpha_{t,0} + 1)(\alpha_{t,0} + 2)}$$

# Algorithm

## Algorithm1 : Posterior Update

---

### Algorithm 1 PosteriorUpdate( $\alpha_t, \mathbf{e}$ )

---

```
1: for  $i = 1, 2 \dots N - 1$  do                                ▷ Compute posterior moments
2:    $m_i \leftarrow \frac{\alpha_{t,i}(e_i + \mathbf{e} \cdot \alpha_t)}{(\mathbf{e} \cdot \alpha_t)(\alpha_{t,0} + 1)}$ 
3:  $s_1 \leftarrow \frac{\alpha_{t,1}(\alpha_{t,1} + 1)(2e_1 + \mathbf{e} \cdot \alpha_t)}{(\mathbf{e} \cdot \alpha_t)(\alpha_{t,0} + 1)(\alpha_{t,0} + 2)}$ 
4:  $\alpha_{t+1,0} \leftarrow \frac{m_1 - s_1}{s_1 - m_1}$                                 ▷ Compute  $\alpha_{t+1}$ 
5: for  $i = 1, 2 \dots N - 1$  do
6:    $\alpha_{t+1,i} \leftarrow m_i \alpha_{t+1,0}$ 
7:  $\alpha_{t+1,N} \leftarrow \alpha_{t+1,0} - \sum_{i=1}^{N-1} \alpha_{t+1,i}$ 
8: return  $\alpha_{t+1}$ 
```

---

## Algorithm 2 : RL with Bayesian Reward Shaping

---

### Algorithm 2 RL with Bayesian Reward Shaping

---


```
1: initialize  $\alpha \in \mathbb{R}_+^N$ 
2: for  $episode = 0, 1 \dots M$  do                                ▷ Main loop
3:    $\hat{\Phi} \leftarrow \frac{\sum_{i=1}^N \Phi_i \alpha_i}{\sum_{i=1}^N \alpha_i}$                                 ▷ Pool experts and compute shaped reward
4:    $F(s, a, s') \leftarrow \gamma \hat{\Phi}(s') - \hat{\Phi}(s)$ 
5:    $(R_t, s_t)_{t=1 \dots T} \leftarrow \text{TrainRL}(F)$                                 ▷ Perform one episode of training
6:   for all  $(R_t, s_t)$  do                                ▷ Posterior update
7:     update  $\hat{\sigma}^2$  and compute  $\mathbf{e}$ 
8:      $\alpha \leftarrow \text{PosteriorUpdate}(\alpha, \mathbf{e})$ 
```

---



*04.*

## Experimental Results



# *CartPole*

This is a classical control problem implemented in OpenAI Gym


Parameters	Value
$\gamma$	0.95
$\varepsilon$	1, Decay:0.98
$\alpha$	0.5, Decay:0.99
Reward	+1 at every step as long as the pole is upright
Stoppage	Last 5 episodes is 500

The expert used :


- Optimal: Complete trained agent
- Good : Partially trained function
- Bad : -good
- zero: Return 0 only



# CartPole



Q learning



SARSA

